

## First lets read in the data

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sb
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.tree import DecisionTreeClassifier
from sklearn import preprocessing
from sklearn.metrics import classification_report
df = pd.read_csv("Auto.csv")
print(df.head())
print('\nDimensions of data frame:', df.shape)
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70.0	
1	15.0	8	350.0	165	3693	11.5	70.0	
2	18.0	8	318.0	150	3436	11.0	70.0	
3	16.0	8	304.0	150	3433	12.0	70.0	
4	17.0	8	302.0	140	3449	NaN	70.0	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

Dimensions of data frame: (392, 9)

## Data exploration

```
In [2]: df[['mpg', 'weight', 'year']].describe()
```

```
Out[2]:
```

	mpg	weight	year
<b>count</b>	392.000000	392.000000	390.000000
<b>mean</b>	23.445918	2977.584184	76.010256
<b>std</b>	7.805007	849.402560	3.668093
<b>min</b>	9.000000	1613.000000	70.000000
<b>25%</b>	17.000000	2225.250000	73.000000
<b>50%</b>	22.750000	2803.500000	76.000000
<b>75%</b>	29.000000	3614.750000	79.000000
<b>max</b>	46.600000	5140.000000	82.000000

The range of the columns in order are as follows (37.6, 3527, 10). The average of each of the columns in order are as follows (23.445918, 2977.584184, 76.010256).

---

## Now we're going to do some more data exploration

---

```
In [3]: df.dtypes
```

```
Out[3]: mpg           float64
cylinders          int64
displacement       float64
horsepower         int64
weight            int64
acceleration       float64
year              float64
origin            int64
name              object
dtype: object
```

## Lets change some of these types and output them again

---

```
In [4]: df.cylinders = df.cylinders.astype('category').cat.codes
df['origin'] = df['origin'].astype('category')
df.dtypes
```

```
Out[4]: mpg           float64
cylinders          int8
displacement       float64
horsepower         int64
weight            int64
acceleration       float64
year              float64
origin            category
name              object
dtype: object
```

I tried using `df.cylinders = df.cylinders.astype('category').cat.codes` to change the first one, but for some reason it would not change. I resigned to just using the same method for both as to continue with the assignment.

---

## Now lets get rid of the rows with NA

---

```
In [5]: df.isnull().sum()
```

```
Out[5]: mpg          0
        cylinders    0
        displacement  0
        horsepower    0
        weight        0
        acceleration  1
        year          2
        origin        0
        name          0
        dtype: int64
```

```
In [6]: df = df.dropna()
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: mpg          0
        cylinders    0
        displacement  0
        horsepower    0
        weight        0
        acceleration  0
        year          0
        origin        0
        name          0
        dtype: int64
```

```
In [8]: print('\nDimensions of data frame:', df.shape)
```

Dimensions of data frame: (389, 9)

## Lets add the mpg\_high column

```
In [9]: df['mpg_high'] = np.where(df['mpg'] > df['mpg'].mean(), 1, 0)
        df['mpg_high'] = df['mpg_high'].astype('category')
```

## Now lets remove the mpg column as well as the name column

```
In [10]: df = df.drop(columns = ['mpg', 'name'])
         df.head()
```

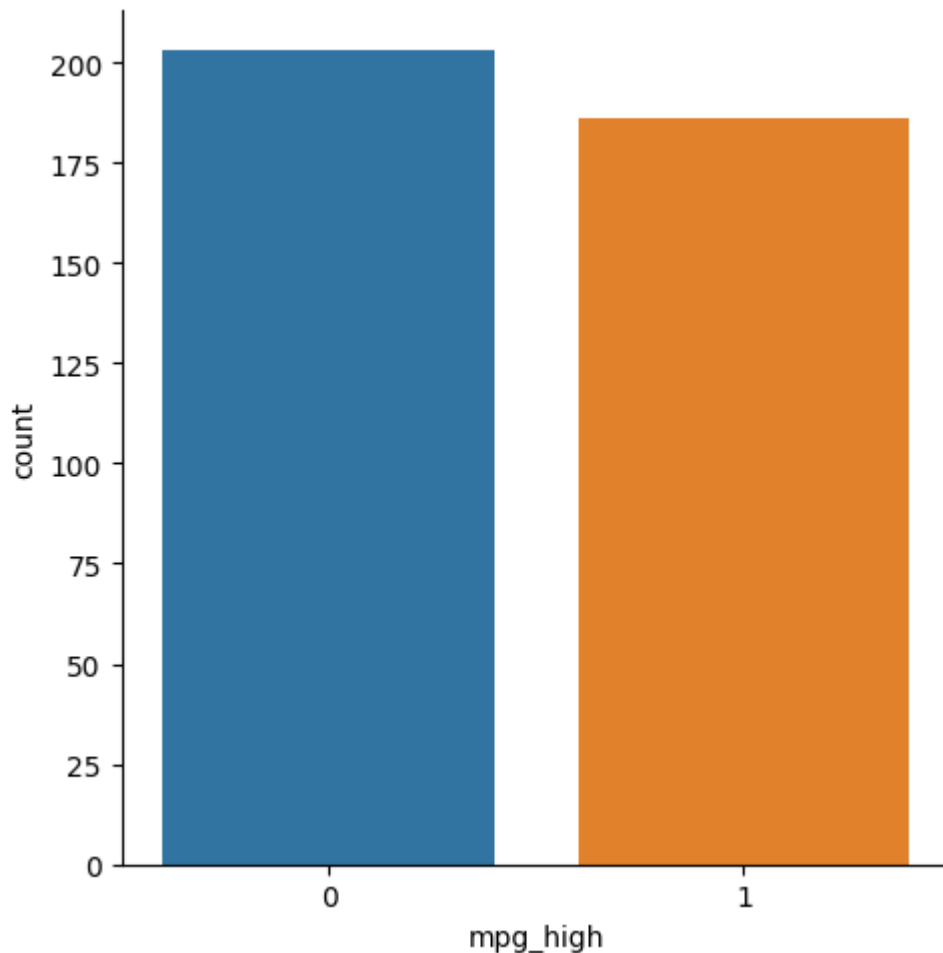
```
Out[10]:
```

	cylinders	displacement	horsepower	weight	acceleration	year	origin	mpg_high
0	4	307.0	130	3504	12.0	70.0	1	0
1	4	350.0	165	3693	11.5	70.0	1	0
2	4	318.0	150	3436	11.0	70.0	1	0
3	4	304.0	150	3433	12.0	70.0	1	0
6	4	454.0	220	4354	9.0	70.0	1	0

## Data Exploration with graphs

```
In [11]: sb.catplot(x = "mpg_high", kind = "count", data = df)
```

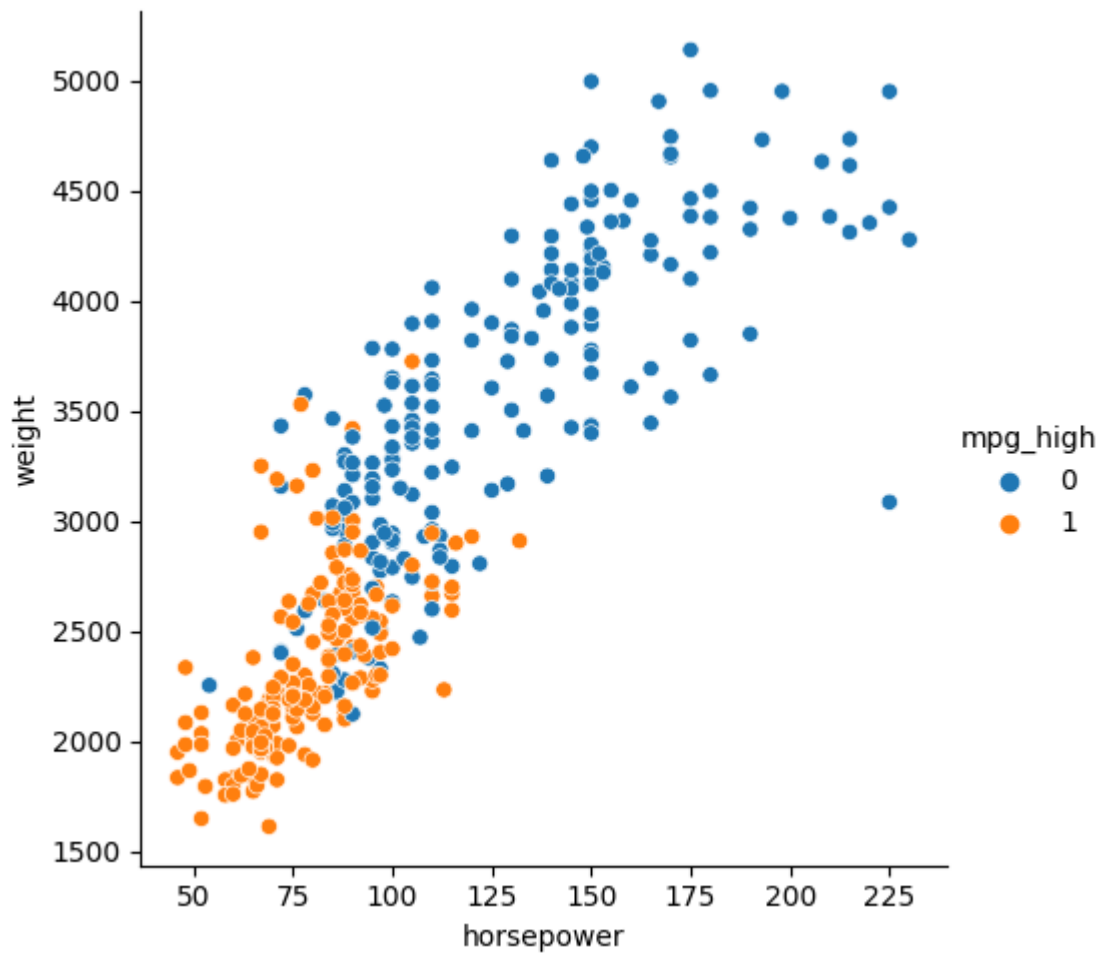
```
Out[11]: <seaborn.axisgrid.FacetGrid at 0x1f2bfa363b0>
```



Using this graph we can see that there are slightly more cars that have a less than average mpg, which means that the cars with higher than average mpg near the top must be a bit higher than the cars with lower than average mpg near the bottom.

```
In [12]: sb.relplot(x = "horsepower", y = "weight", data = df, hue = "mpg_high")
```

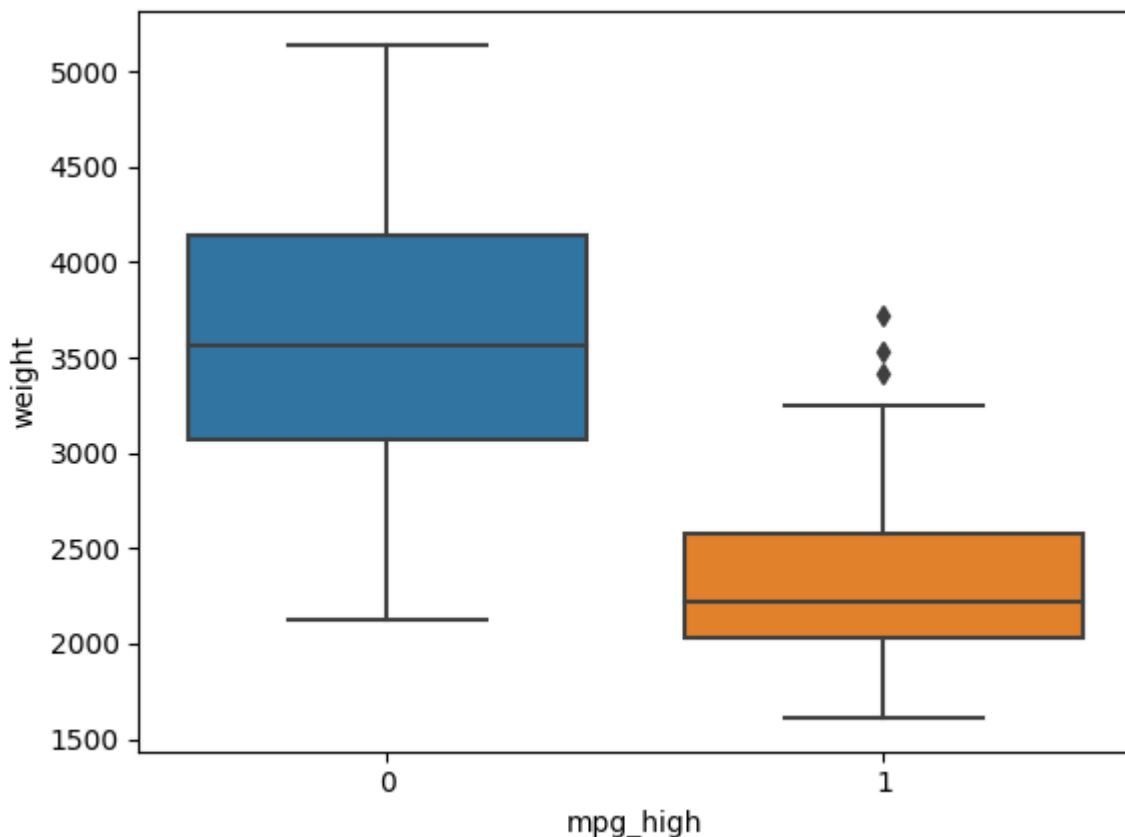
```
Out[12]: <seaborn.axisgrid.FacetGrid at 0x1f2bfa85ed0>
```



This graph tells us that heavier vehicles tend to have higher horsepower and also consume more fuel.

```
In [13]: sb.boxplot(x = "mpg_high", y = "weight", data = df)
```

```
Out[13]: <AxesSubplot: xlabel='mpg_high', ylabel='weight'>
```



This graph tells us that of the vehicles that use more gas, they are heavier and the lower boundary for these vehicles weight is around the average weight of a lighter vehicle. Also the vehicles that weigh more have a higher range of weights compared to those that weigh less and consume less gas. There are also a couple of outliers in the vehicles that use less gas that are heavier than the others. These are probably SUV's where as the heavier vehicles are probably trucks.

## Lets divide the data into test and train

```
In [14]: X = df.loc[:, df.columns != 'mpg_high']
y = df.mpg_high
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2, random_state =
print('\nDimensions of train data frame:', X_train.shape)
print('\nDimensions of test frame:', X_test.shape)
```

Dimensions of train data frame: (311, 7)

Dimensions of test frame: (78, 7)

## Now lets train a Logistic Regression model

```
In [15]: clf = LogisticRegression(max_iter = 400, random_state = 1234)
clf.fit(X_train, y_train)
clf.score(X_train, y_train)
pred = clf.predict(X_test)
```

```
In [16]: print('accuracy score: ', accuracy_score(y_test, pred))
print('precision score: ', precision_score(y_test, pred))
print('recall score: ', recall_score(y_test, pred))
print('f1 score: ', f1_score(y_test, pred))
```

```
accuracy score:  0.8974358974358975
precision score:  0.7777777777777778
recall score:    1.0
f1 score:        0.8750000000000001
```

It seems that using Logistic Regression we can get a pretty good idea of if a car will consume a lot of gas or not based on the other factors provided such as weight, horsepower, year, etc.

## Now lets try a Decision Tree

```
In [17]: clf = DecisionTreeClassifier(random_state = 1234)
clf.fit(X_train, y_train)
pred = clf.predict(X_test)
```

```
In [18]: print('accuracy score: ', accuracy_score(y_test, pred))
print('precision score: ', precision_score(y_test, pred))
print('recall score: ', recall_score(y_test, pred))
print('f1 score: ', f1_score(y_test, pred))
```

```
accuracy score:  0.9230769230769231
precision score:  0.8666666666666667
recall score:    0.9285714285714286
f1 score:        0.896551724137931
```

here the decision tree got a better accuracy and a much better precision score but about 8 percent less on the recall score. The precision score makes sense to me since DT puts the data points into sort of "buckets" so they would naturally be close together compared to Logistic Regression in some cases.

## Now lets do a Neural Network

First lets normalize the data.

```
In [19]: scaler = preprocessing.StandardScaler().fit(X_train)

X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## Now lets train the NN

```
In [20]: clf = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(5, 2), max_iter=1000, random_s
clf.fit(X_train_scaled, y_train)
pred = clf.predict(X_test_scaled)
```

## Lets see how the NN did

```
In [21]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.93	0.86	0.90	50
1	0.78	0.89	0.83	28
accuracy			0.87	78
macro avg	0.86	0.88	0.86	78
weighted avg	0.88	0.87	0.87	78

The NN here actually did slightly worse than Logistic Regression, let's try with some different settings and see how this changes.

```
In [22]: clf = MLPClassifier(solver='adam', hidden_layer_sizes=(100,), max_iter=2000, random_state=42)
clf.fit(X_train_scaled, y_train)
pred = clf.predict(X_test_scaled)
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.93	0.86	0.90	50
1	0.78	0.89	0.83	28
accuracy			0.87	78
macro avg	0.86	0.88	0.86	78
weighted avg	0.88	0.87	0.87	78

It seems that changing some of the settings had little to no difference on the outcome.

```
In [23]: clf = MLPClassifier(solver='sgd', hidden_layer_sizes=(100,), max_iter=1000, random_state=42)
clf.fit(X_train_scaled, y_train)
pred = clf.predict(X_test_scaled)
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.98	0.82	0.89	50
1	0.75	0.96	0.84	28
accuracy			0.87	78
macro avg	0.86	0.89	0.87	78
weighted avg	0.89	0.87	0.87	78

## Analysis

Out of the models that we attempted here the best results by about 4 percent was the Decision Tree. I partially expected better outcomes from the NN, perhaps they would perform better with some other combination of settings. The DT doing better than LR is no surprise, as it seemed from the graphs that the data was grouped quite well on one axis but maybe not the others. The recall score of LR makes me think that it may have overfitted a bit or perhaps the outliers



negatively affected it more. I personally really liked using sklearn as opposed to R. I think that here you can sort of feel that it was written for Computer Scientists where R feels that it was made for Data Scientists or Statisticians. Especially when it comes to manipulating the data, I thought that writing the python code felt much more natural.