

# Classification

Here is the data used for classification.

---

## Dividing Data

```
library(readr)
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v dplyr   1.0.10
## v tibble  3.1.8      v stringr 1.4.1
## v tidyr   1.2.1      v forcats 0.5.2
## v purrr   0.3.4
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(dplyr)
library(ROCR)
library(xgboost)

##
## Attaching package: 'xgboost'
##
## The following object is masked from 'package:dplyr':
##
##   slice

library(mccr)
library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:dplyr':
##
##   combine
##
## The following object is masked from 'package:ggplot2':
##
##   margin

library(ISLR)
library(caret)

## Loading required package: lattice
```

```
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##   lift
library(adabag)

## Loading required package: rpart
## Loading required package: foreach
##
## Attaching package: 'foreach'
##
## The following objects are masked from 'package:purrr':
##
##   accumulate, when
##
## Loading required package: doParallel
## Loading required package: iterators
## Loading required package: parallel

library(tree)
library(rpart)
library(e1071)
library(tree)

df <- read.csv("credit.csv")

set.seed(1234)
i <- sample(1:nrow(df), nrow(df)*0.8, replace=FALSE)
train <- df[i,]
test <- df[-i,]

set.seed(1234)
train_2 <- train
train_2$dpm <- as.factor(train_2$dpm)
```


---

## Data exploration


Lets take a look at the data graphically as well as some of the statistics.

```
hist(df$LIMIT_BAL)
```


**Histogram of df\$LIMIT\_BAL**



```
plot(df$AGE, df$LIMIT_BAL)
```



```
plot(df$EDUCATION, df$LIMIT_BAL)
```



```
cor(df[2:6], use = "complete")
```


```
##          LIMIT_BAL          SEX  EDUCATION  MARRIAGE          AGE
## LIMIT_BAL  1.00000000  0.02475524 -0.21916070 -0.10813941  0.14471280
## SEX        0.02475524  1.00000000  0.01423194 -0.03138884 -0.09087365
## EDUCATION -0.21916070  0.01423194  1.00000000 -0.14346434  0.17506066
## MARRIAGE  -0.10813941 -0.03138884 -0.14346434  1.00000000 -0.41416992
## AGE       0.14471280 -0.09087365  0.17506066 -0.41416992  1.00000000
```

## Decision Tree Baseline

```
start_time <- Sys.time()
tree1 <- tree(as.factor(dpnm) ~., data=train)
end_time <- Sys.time()
summary(tree1)
```

```
##
## Classification tree:
## tree(formula = as.factor(dpnm) ~ ., data = train)
## Variables actually used in tree construction:
## [1] "PAY_1" "PAY_2" "PAY_AMT3"
## Number of terminal nodes: 4
## Residual mean deviance: 0.8926 = 21420 / 24000
## Misclassification error rate: 0.1809 = 4341 / 24000
```

```
plot(tree1)
text(tree1, cex=1)
```



```
print(paste("Training time: ", (end_time-start_time), "s"))
```

```
## [1] "Training time: 0.296769142150879 s"
```

```
pred <- predict(tree1,newdata=test,type="class")
table(pred,test$dpnm)
```

```
##
## pred    0    1
##      0 4526  897
##      1  174  403
```

```
acc <- mean(pred==test$dpnm)
mcc <- mCCR(factor(pred), test$dpnm)
print(paste("accuracy=", acc))
```

```
## [1] "accuracy= 0.8215"
```

```
print(paste("mcc=", mcc))
```

```
## [1] "mcc= 0.381453579643801"
```

```
##Random Forest
```

```
start_time <- Sys.time()
rf <- randomForest(dpnm ~ ., data=train_2, importance = TRUE, proximity = TRUE)
end_time <- Sys.time()
```

```
summary(rf)
```

```
##           Length   Class Mode
## call                5 -none- call
## type                1 -none- character
## predicted          24000 factor numeric
## err.rate           1500 -none- numeric
## confusion           6 -none- numeric
## votes              48000 matrix numeric
## oob.times           24000 -none- numeric
## classes             2 -none- character
## importance           96 -none- numeric
## importanceSD         72 -none- numeric
## localImportance       0 -none- NULL
## proximity           576000000 -none- numeric
## ntree                1 -none- numeric
## mtry                 1 -none- numeric
## forest              14 -none- list
## y                   24000 factor numeric
## test                0 -none- NULL
## inbag               0 -none- NULL
## terms               3 terms  call

print(paste("Training time: ", (end_time-start_time), "s"))
```

```
## [1] "Training time: 7.65738598108292 s"
```

```
pred <- predict(rf, newdata = test)
acc <- mean(pred == test$dpm)
mcc <- mccr(factor(pred), test$dpm)
table(pred, test$dpm)
```

```
##
## pred    0    1
##    0 4451  840
##    1  249  460
```

```
print(paste("Accuracy = ", acc))
```

```
## [1] "Accuracy = 0.8185"
```

```
print(paste("MCC = ", mcc))
```

```
## [1] "MCC = 0.383975737946668"
```

## XGBoost

```
start_time <- Sys.time()
xg <- xgboost(data = data.matrix(train), label = train$dpm, nrounds = 100)
```

```
## [1] train-rmse:0.350012
## [2] train-rmse:0.245017
## [3] train-rmse:0.171518
## [4] train-rmse:0.120067
## [5] train-rmse:0.084050
## [6] train-rmse:0.058837
## [7] train-rmse:0.041187
## [8] train-rmse:0.028832
## [9] train-rmse:0.020183
```

```
## [10] train-rmse:0.014129
## [11] train-rmse:0.009891
## [12] train-rmse:0.006924
## [13] train-rmse:0.004847
## [14] train-rmse:0.003393
## [15] train-rmse:0.002375
## [16] train-rmse:0.001663
## [17] train-rmse:0.001164
## [18] train-rmse:0.000815
## [19] train-rmse:0.000570
## [20] train-rmse:0.000399
## [21] train-rmse:0.000279
## [22] train-rmse:0.000196
## [23] train-rmse:0.000137
## [24] train-rmse:0.000096
## [25] train-rmse:0.000067
## [26] train-rmse:0.000047
## [27] train-rmse:0.000033
## [28] train-rmse:0.000023
## [29] train-rmse:0.000016
## [30] train-rmse:0.000011
## [31] train-rmse:0.000008
## [32] train-rmse:0.000006
## [33] train-rmse:0.000005
## [34] train-rmse:0.000005
## [35] train-rmse:0.000005
## [36] train-rmse:0.000005
## [37] train-rmse:0.000005
## [38] train-rmse:0.000005
## [39] train-rmse:0.000005
## [40] train-rmse:0.000005
## [41] train-rmse:0.000005
## [42] train-rmse:0.000005
## [43] train-rmse:0.000005
## [44] train-rmse:0.000005
## [45] train-rmse:0.000005
## [46] train-rmse:0.000005
## [47] train-rmse:0.000005
## [48] train-rmse:0.000005
## [49] train-rmse:0.000005
## [50] train-rmse:0.000005
## [51] train-rmse:0.000005
## [52] train-rmse:0.000005
## [53] train-rmse:0.000005
## [54] train-rmse:0.000005
## [55] train-rmse:0.000005
## [56] train-rmse:0.000005
## [57] train-rmse:0.000005
## [58] train-rmse:0.000005
## [59] train-rmse:0.000005
## [60] train-rmse:0.000005
## [61] train-rmse:0.000005
## [62] train-rmse:0.000005
## [63] train-rmse:0.000005
```



```
## [64] train-rmse:0.000005
## [65] train-rmse:0.000005
## [66] train-rmse:0.000005
## [67] train-rmse:0.000005
## [68] train-rmse:0.000005
## [69] train-rmse:0.000005
## [70] train-rmse:0.000005
## [71] train-rmse:0.000005
## [72] train-rmse:0.000005
## [73] train-rmse:0.000005
## [74] train-rmse:0.000005
## [75] train-rmse:0.000005
## [76] train-rmse:0.000005
## [77] train-rmse:0.000005
## [78] train-rmse:0.000005
## [79] train-rmse:0.000005
## [80] train-rmse:0.000005
## [81] train-rmse:0.000005
## [82] train-rmse:0.000005
## [83] train-rmse:0.000005
## [84] train-rmse:0.000005
## [85] train-rmse:0.000005
## [86] train-rmse:0.000005
## [87] train-rmse:0.000005
## [88] train-rmse:0.000005
## [89] train-rmse:0.000005
## [90] train-rmse:0.000005
## [91] train-rmse:0.000005
## [92] train-rmse:0.000005
## [93] train-rmse:0.000005
## [94] train-rmse:0.000005
## [95] train-rmse:0.000005
## [96] train-rmse:0.000005
## [97] train-rmse:0.000005
## [98] train-rmse:0.000005
## [99] train-rmse:0.000005
## [100] train-rmse:0.000005
```

```
end_time <- Sys.time()
summary(xg)
```

```
##           Length Class           Mode
## handle           1 xgb.Booster.handle externalptr
## raw             75159 -none-         raw
## niter            1 -none-         numeric
## evaluation_log    2 data.table       list
## call             13 -none-         call
## params            1 -none-         list
## callbacks          2 -none-         list
## feature_names     25 -none-         character
## nfeatures          1 -none-         numeric
```

```
print(paste("Training time: ", (end_time-start_time), "s"))
```

```
## [1] "Training time: 0.669389963150024 s"
```

```

pred <- predict(xg, data.matrix(test), reshape=TRUE)

prediction <- as.numeric(pred > 0.5)

err <- mean(prediction != test$dpm)
print(paste("test-error=", err))

## [1] "test-error= 0"

```

## Adabag

```

start_type <- Sys.time()
adab <- boosting(dpm~. , data=train_2, boos = TRUE, mfinal=10)
end_time <- Sys.time()
summary(adab)

```

```

##           Length Class  Mode
## formula         3 formula call
## trees           10 -none- list
## weights         10 -none- numeric
## votes          48000 -none- numeric
## prob            48000 -none- numeric
## class           24000 -none- character
## importance       24 -none- numeric
## terms            3 terms call
## call             5 -none- call

```

```

print(paste("Training time: ", (end_time-start_time), "s"))

```

```

## [1] "Training time:  20.3851799964905 s"

```

```

pred <- predict(adab, newdata=test, type="response")
pred$class <- as.integer(pred$class)
acc <- mean(pred$class==test$dpm)
print(paste("accuracy=", acc))

```

```

## [1] "accuracy= 0.817166666666667"

```

##Analysis

Decision Tree - As a baseline, the decision tree provides an accuracy of 82.15%, with an MCC of 0.38, making it fairly accurate, considering that it only took 0.33s to compute, making it nearly instantaneous.

Random Forest - The random Forest algorithm, on the other hand has an accuracy of 81.9%, and an mcc of 0.38, but takes over 10x as long, taking a little over 7 minutes to fully create a model. Because there is not even an increase in accuracy in exchange for the computation time.

XGBoost has a accuracy that is around the same as Random Forest, but is able to do it in a fraction of the time, even coming in faster than the Decision Tree algorithm.

Adaboost is the last algorithm, and this completels in 2.65s with an accuracy of 82.3%, which is not significantly higher than any of the others.

Conclusion - While all the models, and methods were able to maintain a similar level of accuracy while predicting, the biggest difference between them was the time that it took to compute. From all the results seen, it seems that XGBoost is the best technique to use for this type of machine learning. Another option for fast binary classification would be Fast Adaboost.