

SVM Regression

I decided to use the bmw used car dataset again for regression as it has provided good results in the past.
<https://www.kaggle.com/datasets/adityadesail3/used-car-dataset-ford-and-mercedes>

Lets go ahead and load in the data, divide it into train and test data, and factorize it.

```
library(readr)
df <- read.csv("bmw.csv")

df$model <- as.factor(df$model)
df$transmission <- as.factor(df$transmission)
df$fuelType <- as.factor(df$fuelType)


set.seed(1)
groups <- c(train=.6, test=.2, validate=.2)
i <- sample(cut(1:nrow(df),nrow(df)*cumsum(c(0,groups))), labels=names(groups))
train <- df[i=="train",]
test <- df[i=="test",]
vald <- df[i=="validate",]
```

Lets take a look at some of the data.


```
library(ggplot2)
head(train)
```

##		model	year	price	transmission	mileage	fuelType	tax	mpg	engineSize
## 1	5	Series	2014	11200	Automatic	67068	Diesel	125	57.6	2.0
## 3	5	Series	2016	16000	Automatic	62794	Diesel	160	51.4	3.0
## 8	2	Series	2018	16250	Manual	10401	Petrol	145	52.3	1.5
## 10	5	Series	2016	14250	Automatic	36099	Diesel	20	68.9	2.0
## 12	1	Series	2017	11800	Manual	29840	Diesel	20	68.9	2.0
## 15		X3	2017	22000	Automatic	19057	Diesel	145	54.3	2.0

```
ggplot(train, aes(x = price)) + geom_histogram(fill="darkgreen", binwidth = 1000)
```



```
ggplot(train, aes(x = price, y = mileage)) + geom_point(aes(color = transmission))
```



Interestingly it seems that there might be some correlation between the type of transmission and the price. It seems that the manual cars tend to have lower prices despite maybe having lower mileage as well, where low mileage usually means higher price. This is followed by the automatic transmissions having the next higher price and the semi-auto with the best price.

Lets start with linear regression

First we will tune the hyper parameter C to find the best possible value.

```
library(e1071)
tune.out <- tune(svm, price~., data=vald, kernel="linear",
ranges=list(cost=c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 23630283
##
## - Detailed performance results:
```

```
##      cost      error dispersion
## 1 1e-03 61449437   17919439
## 2 1e-02 41796289   15589739
## 3 1e-01 28223883   15341158
## 4 1e+00 23630283   16805083
## 5 5e+00 23793358   17092360
## 6 1e+01 23810869   17120243
## 7 1e+02 23833289   17179911
```

It seems the best c value is 1, so lets use it for the actual model now.

```
svm_linear <- svm(price~., data = train, kernel = "linear", cost = 1, scale = TRUE)
pred_linear <- predict(svm_linear, newdata = test)
cor(pred_linear, test$price)
```

```
## [1] 0.9338052
```

Looks like linear gets a correlation value of .933 which is quite good.

Now lets try out polynomial

We can tune the C hyper parameter the same way but for polynomial.

```
library(e1071)
tune.out <- tune(svm, price~., data=vald, kernel="polynomial",
ranges=list(cost=c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##      cost
##        1
##
## - best performance: 95684474
##
## - Detailed performance results:
##      cost      error dispersion
## 1 1e-03 144148052   29988639
## 2 1e-02 140100022   29874489
## 3 1e-01 123973024   29270559
## 4 1e+00  95684474   29158914
## 5 5e+00 118598149  180686584
## 6 1e+01 159636339  334379429
## 7 1e+02 139099067  295603586
```

Polynomial had a best value of 100, so lets train the model using 100.

```
svm_poly <- svm(price~., data = train, kernel = "polynomial", cost = 100, scale = TRUE)
pred_poly <- predict(svm_poly, newdata = test)
cor(pred_poly, test$price)
```

```
## [1] 0.9638701
```

Polynomial gets a correlation value of .963 which is even better than linear and an amazing value.

Last we try radial

First we need to tune the C hyper parameter like before, but now we also need to tune gamma.

```
tune.out <- tune(svm, price~., data=vald, kernel="radial",
ranges=list(cost=c(0.001, 0.01, 0.1, 1, 5, 10, 100),
            gamma=c(.5,1,2,3,4)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##     5    0.5
##
## - best performance: 18419729
##
## - Detailed performance results:
##   cost gamma   error dispersion
## 1  1e-03   0.5 134365501  31075114
## 2  1e-02   0.5 88442965  29681584
## 3  1e-01   0.5 37556984  23420182
## 4  1e+00   0.5 19392539  18654620
## 5  5e+00   0.5 18419729  17983828
## 6  1e+01   0.5 18797604  17965031
## 7  1e+02   0.5 21448376  17750152
## 8  1e-03   1.0 140818560  31084772
## 9  1e-02   1.0 116106601  30929869
## 10 1e-01   1.0 53945232  26169102
## 11 1e+00   1.0 23781780  18881043
## 12 5e+00   1.0 22395267  18181565
## 13 1e+01   1.0 22647633  18216508
## 14 1e+02   1.0 25367115  17571796
## 15 1e-03   2.0 142954002  31077677
## 16 1e-02   2.0 130538642  30929884
## 17 1e-01   2.0 81543141  29843770
## 18 1e+00   2.0 34482454  19549585
## 19 5e+00   2.0 29964181  18804904
## 20 1e+01   2.0 30358155  18773531
## 21 1e+02   2.0 33704805  18390627
## 22 1e-03   3.0 143330741  31073156
## 23 1e-02   3.0 133624390  30835807
## 24 1e-01   3.0 90942722  30386154
## 25 1e+00   3.0 42296461  20629327
## 26 5e+00   3.0 35302796  19287202
## 27 1e+01   3.0 35767988  19275202
## 28 1e+02   3.0 39836936  19278808
## 29 1e-03   4.0 143486421  31067904
```

```
## 30 1e-02    4.0 134918480    30849986
## 31 1e-01    4.0  95461005    30389511
## 32 1e+00    4.0 47004767    21259124
## 33 5e+00    4.0 39074511    19496033
## 34 1e+01    4.0 39417619    19464308
## 35 1e+02    4.0 43193192    19428237
```

The best value for cost is 5 and gamma is 0.5, so we will use these parameters for the model.

```
svm_rad <- svm(price~., data = train, kernel = "radial", cost = 5, gamma = 0.5, scale = TRUE)
pred_rad <- predict(svm_rad, newdata = test)
cor(pred_rad, test$price)
```

```
## [1] 0.9777572
```

Radial got a correlation value of roughly .978 which is the best so far and overall an amazing predictor.

Analysis

The first thing that should be noted is that all three of the kernel methods got amazing correlation values. However, as we moved to higher dimensions the correlation value increased each time. The first increase from linear to polynomial was the biggest by about 3 percent and then another increase of about 1.5 percent from polynomial to radial. The fact that polynomial did better than linear suggests that the data can be better split when considering more dimensions to the data. Using the polynomial kernel methods we were able to improve, and with radial we are able to control the variance and bias of the fitting. By having the model use a low gamma value it probably had higher bias and lower variance, which also improved the model slightly.